

Figure 1 Learned gate per layer and head (n = 16 runs)

